

$$p_{\theta_{dec}}(\mathbf{y}_{1:m}|\mathbf{c}).$$

By Bayes' rule the distribution can be decomposed into conditional distributions of single target vectors as follows:

$$p_{\theta_{dec}}(\mathbf{Y}_{1:m}|\mathbf{c}) = \prod_{i=1}^m p_{\theta_{dec}}(\mathbf{y}_i|\mathbf{Y}_{0:i-1}, \mathbf{c}).$$

Thus, if the architecture can model the conditional distribution of the next target vector, given all previous target vectors:

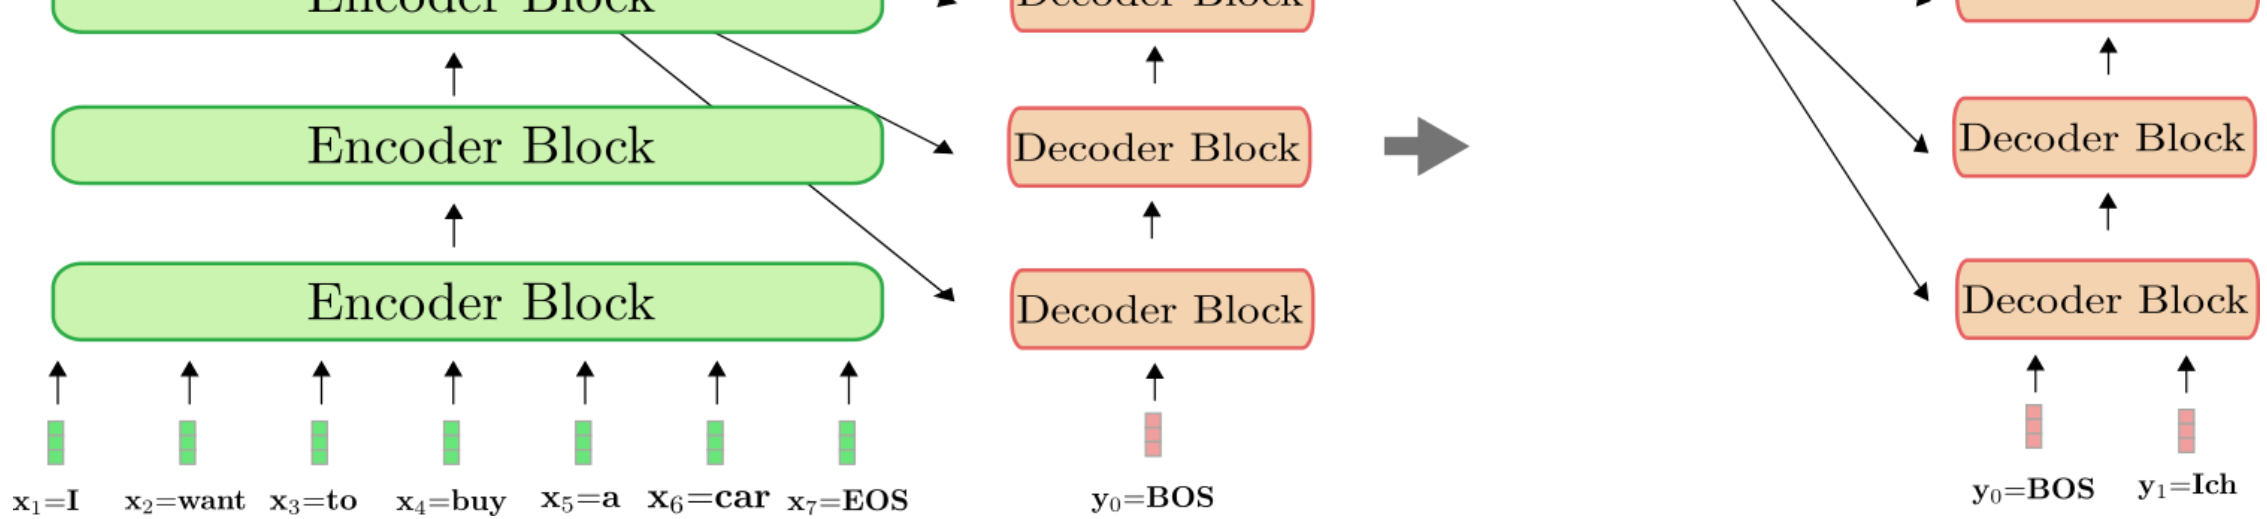
$$p_{\theta_{dec}}(\mathbf{y}_i|\mathbf{Y}_{0:i-1}, \mathbf{c}), \forall i \in \{1, \dots, m\},$$

then it can model the distribution of any target vector sequence given the hidden state  $\mathbf{c}$  by simply multiplying all conditional probabilities.

So how does the RNN-based decoder architecture model  $p_{\theta_{dec}}(\mathbf{y}_i|\mathbf{Y}_{0:i-1}, \mathbf{c})$ ?

In computational terms, the model sequentially maps the previous inner hidden state  $\mathbf{c}_{i-1}$  and the previous target vector  $\mathbf{y}_i$  to the current inner hidden state  $\mathbf{c}_i$  and a *logit vector*  $\mathbf{l}_i$  (shown in dark red below):

$$f_{\theta_{dec}}(\mathbf{y}_{i-1}, \mathbf{c}_{i-1}) \rightarrow \mathbf{l}_i, \mathbf{c}_i$$



As can be seen, only in step  $i = 1$  do we have to encode "I want to buy a car EOS" to  $\overline{\mathbf{X}}_{1:7}$ . At step  $i = 2$ , the contextualized encodings of "I want to buy a car EOS" are simply reused by the decoder.

In 🤗 Transformers, this auto-regressive generation is done under-the-hood when calling the `.generate()` method. Let's use one of our translation models to see this in action.

## Encoder

As mentioned in the previous section, the *transformer-based* encoder maps the input

directional self-attention.

As in bi-directional self-attention, in uni-directional self-attention, the query vectors  $\mathbf{q}_0, \dots, \mathbf{q}_{m-1}$  (shown in purple below), key vectors  $\mathbf{k}_0, \dots, \mathbf{k}_{m-1}$  (shown in orange below), and value vectors  $\mathbf{v}_0, \dots, \mathbf{v}_{m-1}$  (shown in blue below) are projected from their respective input vectors  $\mathbf{y}'_0, \dots, \mathbf{y}'_{m-1}$  (shown in light red below). However, in uni-directional self-attention, each query vector  $\mathbf{q}_i$  is compared *only* to its respective key vector and all previous ones, namely  $\mathbf{k}_0, \dots, \mathbf{k}_i$  to yield the respective *attention weights*.

This prevents an output vector  $\mathbf{y}''_j$  (shown in dark red below) to include any information about the following input vector  $\mathbf{y}_i$ , with  $i > j$  for all  $j \in \{0, \dots, m-1\}$ . As is the case in bi-directional self-attention, the attention weights are then multiplied by their respective value vectors and summed together.

We can summarize uni-directional self-attention as follows:

$$\mathbf{y}''_i = \mathbf{V}_{0:i} \mathbf{Softmax}(\mathbf{K}_{0:i}^\top \mathbf{q}_i) + \mathbf{y}'_i.$$

Note that the index range of the key and value vectors is  $0 : i$  instead of  $0 : m-1$

<sup>1</sup> The word embedding matrix  $\mathbf{W}_{\text{emb}}$  gives each input word a unique *context-independent* vector representation. This matrix is often fixed as the "LM Head" layer. However, the "LM Head" layer can very well consist of a completely independent "encoded vector-to-logit" weight mapping.

<sup>2</sup> Again, an in-detail explanation of the role the feed-forward layers play in transformer-based models is out-of-scope for this notebook. It is argued in [Yun et. al, \(2017\)](#) that feed-forward layers are crucial to map each contextual vector  $\mathbf{x}'_i$  individually to the desired output space, which the *self-attention* layer does not manage to do on its own. It should be noted here, that each output token  $\mathbf{x}'$  is processed by the same feed-forward layer. For more detail, the reader is advised to read the paper.